

Language Agent Tree Search




One frozen model as a planning agent


Weekend 6 · Frontier Agentic AI

Carlos Cotrini

Where LATS sits



How a model gets to act



- The model still only ever emits **text**. It touches nothing
- The system agrees **in advance** which text counts as a command


Monte Carlo Tree Search, before language models




Monte Carlo Tree Search, before language models




Monte Carlo Tree Search, before language models




Monte Carlo Tree Search, before language models



Monte Carlo Tree Search, before language models



Monte Carlo Tree Search, before language models



The search itself barely changed across twenty years. What changed each time is what supplies the two things it needs: something to propose moves, and something to judge a position. LATS supplies both from one frozen model.

Coulom 2006; Kocsis and Szepesvari 2006; Silver et al., Nature 2016; Silver et al., Science 2018.




Section 1

The LATS method

1


UCB

Exploration



inverse to the
of times
I have tried it

Exploitation




$V(\text{machine})$

$UCB(\text{machine}) = V +$

$$\sqrt{\frac{1}{\text{\#times}}}$$

UCT: the critic, plus an exploration bonus

$$\text{UCT}(s) = \underbrace{V(s)}_{\text{the critic}} + \underbrace{w \sqrt{\frac{\ln N(p)}{N(s)}}}_{\text{the exploration bonus}}$$



UCT: the critic, plus an exploration bonus

$$\text{UCT}(s) = \underbrace{V(s)}_{\text{the critic}} + \underbrace{w \sqrt{\frac{\ln N(p)}{N(s)}}}_{\text{the exploration bonus}}$$

The critic

what the model thinks
this branch is worth

UCT: the critic, plus an exploration bonus

$$\text{UCT}(s) = \underbrace{V(s)}_{\text{the critic}} + \underbrace{w \sqrt{\frac{\ln N(p)}{N(s)}}}_{\text{the exploration bonus}}$$

The critic

what the model thinks
this branch is worth

+

The exploration bonus

how little we have tried it

UCT: the critic, plus an exploration bonus

$$\text{UCT}(s) = \underbrace{V(s)}_{\text{the critic}} + \underbrace{w \sqrt{\frac{\ln N(p)}{N(s)}}}_{\text{the exploration bonus}}$$

The critic

what the model thinks
this branch is worth


+

The exploration bonus


how little we have tried it

This morning's words for the same two terms: how good it looked, plus how little we know.

The same loop as this morning, with four edits



The same loop as this morning, with four edits




1 Selection

by UCT

Game over

The same loop as this morning, with four edits




1 Selection

by UCT

2 Expansion

the actor proposes

The same loop as this morning, with four edits




1 Selection by UCT

2 Expansion the *actor* proposes

3 Evaluation the *critic* scores

The same loop as this morning, with four edits



1 Selection

by UCT

2 Expansion

the *actor* proposes


3 Evaluation

the *critic* scores

4 Simulation

the *best* one, to the end

The same loop as this morning, with four edits



1 Selection by UCT


2 Expansion the *actor* proposes

3 Evaluation the *critic* scores

4 Simulation the *best* one, to the end

5 Backpropagation the *real* result, up

The same loop as this morning, with four edits



1 Selection by UCT

2 Expansion the *actor* proposes


3 Evaluation the *critic* scores

4 Simulation the *best* one, to the end

5 Backpropagation the *real* result, up

6 Reflection the *feedback* model writes it

The same loop as this morning, with four edits



1 Selection by UCT

2 Expansion the *actor* proposes

3 Evaluation the *critic* scores

4 Simulation the *best* one, to the end


5 Backpropagation the *real* result, up

6 Reflection the *feedback* model writes it

grey: UCT, what selection compares.

amber: V from the critic, and all that is new.

The same loop as this morning, with four edits



1 Selection by UCT

2 Expansion the *actor* proposes

3 Evaluation the *critic* scores

4 Simulation the *best* one, to the end

5 Backpropagation the *real* result, up

6 Reflection the *feedback* model writes it

into the next prompt

grey: UCT, what selection compares.


amber: V from the critic, and all that is new.

What the critic actually returns

a state 


$$V(s) = \quad +$$

What the critic actually returns



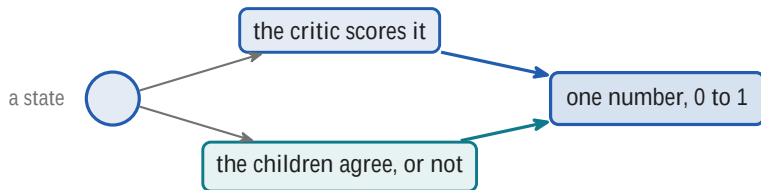
$$V(s) = \quad +$$

What the critic actually returns



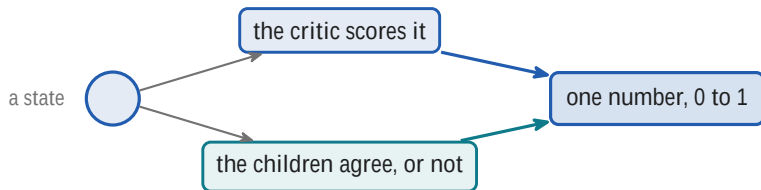
$$V(s) = \quad +$$

What the critic actually returns



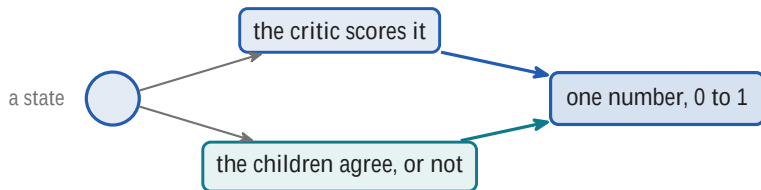
$$V(s) = \quad +$$

What the critic actually returns



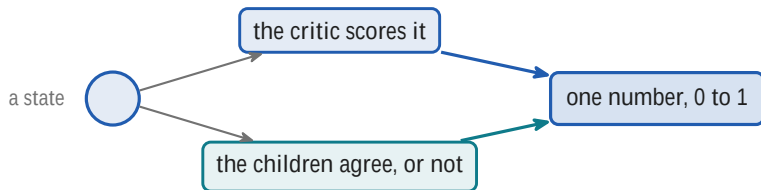
$$V(s) = \lambda LM(s) +$$

What the critic actually returns



$$V(s) = \lambda LM(s) + (1 - \lambda) SC(s)$$

What the critic actually returns




$$V(s) = \lambda LM(s) + (1 - \lambda) SC(s)$$

$\lambda = 0.5$ on question answering, 0.8 on programming and web tasks.


Three jobs, one set of weights

one frozen model


Three jobs, one set of weights




Three jobs, one set of weights



Three jobs, one set of weights



Three jobs, one set of weights



same weights every time. Only the prompt changes.