

FDD Competition Results

Carlos Cotrini

August 28, 2026



Agenda

1. Where it all began: the AML project
2. The challenge: bringing it to you
3. The results: how you really performed
4. The awards: celebrating the top 3

Where it all began

The project has a history

This Kaggle competition was not designed from scratch for you.

- It was the very first graded project in **Advanced Machine Learning (AML)**, an ETH Master's course.
- Every year, Master's students compete against each other on this exact problem, and part of their grade depends on their **position on the leaderboard**.

Why this project?

A different kind of project

- Real, messy data, not a clean notebook exercise.

Why this project?

A different kind of project

- Real, messy data, not a clean notebook exercise.
- Unlike your usual projects, you got **less guidance** and were left more on your own.

Why this project?

A different kind of project

- Real, messy data, not a clean notebook exercise.
- Unlike your usual projects, you got **less guidance** and were left more on your own.
- It's a genuinely **hard** problem: even in AML, Master's students had to compete against each other just to secure their grade on it.

Why this project?

A different kind of project

- Real, messy data, not a clean notebook exercise.
- Unlike your usual projects, you got **less guidance** and were left more on your own.
- It's a genuinely **hard** problem: even in AML, Master's students had to compete against each other just to secure their grade on it.
- There were many possible solutions, and the only way up was to keep improving and **climb the leaderboard** against the other teams.


Why this project?

A different kind of project

- Real, messy data, not a clean notebook exercise.
- Unlike your usual projects, you got **less guidance** and were left more on your own.
- It's a genuinely **hard** problem: even in AML, Master's students had to compete against each other just to secure their grade on it.
- There were many possible solutions, and the only way up was to keep improving and **climb the leaderboard** against the other teams.
- It was a simulation of the real world, where you compete against others on the market for the best model.


AML effort (2025)

- Average **25.6** submissions/team
- Median **19** submissions/team
- Top team: **99** submissions



AML score distribution (2025)

The bar to beat



Mean **0.7115** Median **0.7164** Std **0.0485** Min **0.4865** Max **0.8062**

The challenge

Bringing it to you

So we asked ourselves a question: **how would you perform on a real Machine Learning competition, compared to the AML course?**

- We gave you that very same AML project, with no project structure to lean on.
- The goal was to test your skills on a real problem, on your own.
- You'd compete against each other, exactly like AML students do every year.

The results


And the results were **incredible**.

You **beat** the AML course.


FDD effort

Fewer attempts

- Average **12.6** submissions/team
- Median **5** submissions/team
- One outlier: **96** submissions



FDD score distribution




Mean **0.7426** Median **0.7420** Std **0.0309** Min **0.6774** Max **0.7922**

AML vs FDD, side by side

	AML ($n=136$)	FDD ($n=18$)
Mean	0.7115	0.7426
Median	0.7164	0.7420
Std	0.0485	0.0309
Min	0.4865	0.6774
Max	0.8062	0.7922

Most of you performed above the median



67.6%

Teams of


of AML teams (136 Master's students) scored **at or below** your average score.

1 out of 136

Only **one single AML team** beat your best score.

FDD above AML

Score distributions, normalized to % of teams



Your AML grade

AML grades this project by linear interpolation between 0.5 and 0.8

4.0 6.0

$$\text{grade} = 4 + \frac{\text{score} - 0.5}{0.8 - 0.5} \times 2$$

#	Score	Grade
1	0.7922	5.95
2	0.7892	5.93
3	0.7849	5.90
4	0.7739	5.83
5	0.7644	5.76
6	0.7611	5.74
7	0.7553	5.70
8	0.7454	5.64
9	0.7440	5.63

#	Score	Grade
10	0.7401	5.60
11	0.7351	5.57
12	0.7312	5.54
13	0.7308	5.54
14	0.7188	5.46
15	0.7141	5.43
16	0.7063	5.38
17	0.7032	5.35
18	0.6774	5.18

Every one of you would score above a 5.0.

Congratulations

You beat AML.

- Congratulations to **every** team that submitted: most of you landed above the AML median.
- Everyone should be proud, not just the podium: you took on a genuinely hard, real-world challenge, and performed very well.
- Special congratulations to our **top 3** teams for the podium.

And now...

THE WINNERS

In 3rd place...



R.E.M.

0.7849

In 2nd place...



Team KET

0.7892

And in 1st place...



MagnusAurelius

0.7922

Thank you