



Fine-tuning methods for LLMs

Adapting a model you did not train

Carlos Cotrini



Section 1

From pretraining to adapting



You will start from a released model

Pretraining Apertus-70B

6.30×10^{24} FLOP

- Two days have priced one pretraining run

You will start from a released model

Pretraining Apertus-70B

6.30×10^{24} FLOP

Over 6 million
GPU hours

- Two days have priced one pretraining run

You will start from a released model

Pretraining Apertus-70B


6.30×10^{24} FLOP

Over 6 million
GPU hours

About 4.95 GWh


- Two days have priced one pretraining run

You will start from a released model




- Two days have priced one pretraining run
- None of it is a decision you will ever make

You will start from a released model




- Two days have priced one pretraining run
- None of it is a decision you will ever make
- You cannot pretrain, so you **adapt**, and the same ledger prices that

Sixteen bytes per parameter




Sixteen bytes per parameter



- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$


Sixteen bytes per parameter



- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$
- Apertus-8B needs 128.8 GB of state, and a GH200 holds 96 GB

Sixteen bytes per parameter

14 of the 16 exist only for the parameters you train



= 16 bytes per parameter

- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$
- Apertus-8B needs 128.8 GB of state, and a GH200 holds 96 GB
- Freeze a parameter and 14 of its 16 bytes disappear




Section 2

One matrix, and a low rank update

2

Twenty-four numbers, or ten


W : 6 outputs $\times$ 4 inputs



All 24 are updated

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs



All 24 are updated

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

A : 1×4

a_1	a_2	a_3	a_4
-------	-------	-------	-------

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

A : 1×4

a_1	a_2	a_3	a_4
-------	-------	-------	-------

$b_1 a_1$	$b_1 a_2$	$b_1 a_3$	$b_1 a_4$
$b_2 a_1$	$b_2 a_2$	$b_2 a_3$	$b_2 a_4$
$b_3 a_1$	$b_3 a_2$	$b_3 a_3$	$b_3 a_4$
$b_4 a_1$	$b_4 a_2$	$b_4 a_3$	$b_4 a_4$
$b_5 a_1$	$b_5 a_2$	$b_5 a_3$	$b_5 a_4$
$b_6 a_1$	$b_6 a_2$	$b_6 a_3$	$b_6 a_4$

$$\Delta W = BA$$

- Full fine-tuning updates all 24
- LoRA freezes W and learns B and A : $6 + 4 = 10$ numbers, a ratio of 2.4

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

A : 1×4


a_1	a_2	a_3	a_4
-------	-------	-------	-------

$b_1 a_1$	$b_1 a_2$	$b_1 a_3$	$b_1 a_4$
$b_2 a_1$	$b_2 a_2$	$b_2 a_3$	$b_2 a_4$
$b_3 a_1$	$b_3 a_2$	$b_3 a_3$	$b_3 a_4$
$b_4 a_1$	$b_4 a_2$	$b_4 a_3$	$b_4 a_4$
$b_5 a_1$	$b_5 a_2$	$b_5 a_3$	$b_5 a_4$
$b_6 a_1$	$b_6 a_2$	$b_6 a_3$	$b_6 a_4$

$$\Delta W = BA$$


- Full fine-tuning updates all 24
- LoRA freezes W and learns B and A : $6 + 4 = 10$ numbers, a ratio of 2.4
- Every row of BA is the row A , rescaled. That is what rank one means

The adapter is added, not substituted




$$h = W_0 x$$

The adapter is added, not substituted



$$h = W_0x + BAx$$


The adapter is added, not substituted



$$h = W_0x + BAx$$

- Both paths read the same x , and their outputs are added


The adapter is added, not substituted



$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added


The adapter is added, not substituted



$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added
- α/r rescales the adapter, and it behaves like a learning rate

The adapter is added, not substituted




$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added
- α/r rescales the adapter, and it behaves like a learning rate
- The paper's own tables use ratios of 1, 2 and 8; the PEFT default is $r = 8, \alpha = 8$

B starts at zero, A starts random

$$B = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

B starts at zero, A starts random



B starts at zero, A starts random

$B = 0$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

- The adapter does nothing at all on step one

B starts at zero, A starts random

$B = 0$

0
0
0
0
0
0

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$BA = 0$

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one

B starts at zero, A starts random

$B = 0$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one
- So training starts where pretraining stopped, not somewhere random

B starts at zero, A starts random

$B = 0$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one
- So training starts where pretraining stopped, not somewhere random
- If both started at zero, both gradients would be zero and nothing would ever move


Rank three costs more than the matrix

$r = 1$  10


The matrix itself: 24

⋮


Rank three costs more than the matrix




Rank three costs more than the matrix



Rank three costs more than the matrix




Rank three costs more than the matrix




- Break even at $r = dk/(d + k) = 2.4$ on this matrix

Rank three costs more than the matrix



- Break even at $r = dk/(d + k) = 2.4$ on this matrix
- On a square $d \times d$ matrix that is $r = d/2$, and the saving is $d/(2r)$

Rank three costs more than the matrix



- Break even at $r = dk/(d + k) = 2.4$ on this matrix
- On a square $d \times d$ matrix that is $r = d/2$, and the saving is $d/(2r)$
- Llama-2-7B's query projection is 4096×4096 : at $r = 8$ that is 65 536 numbers against 16 777 216, a factor of 256




Section 3

The same picture at real scale

3

☰ The same picture at real scale

Rank 16, and 0.08 percent of the model trains



8.05×10^9 parameters

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

☰ The same picture at real scale

Rank 16, and 0.08 percent of the model trains



- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
q_proj, v_proj (the PEFT default)	16	6 815 744	0.0846

- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

☰ The same picture at real scale

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
q_proj, v_proj (the PEFT default)	16	6 815 744	0.0846
All six linear layers	16	39 845 888	0.4948

- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
<code>q_proj</code> , <code>v_proj</code> (the PEFT default)	16	6 815 744	0.0846
All six linear layers	16	39 845 888	0.4948
All six linear layers	64	159 383 552	1.9792
All six linear layers	256	637 534 208	7.9167

- The lower bar is drawn to scale, which is why you cannot see it
- Even rank 256 on all six linear layers stays under 8 percent

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state



The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

Weights and gradients	Optimiser state: 96.6 GB
--------------------------	--------------------------


- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB




- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections



16.22 GB, of which 16.11 GB is the frozen weights


- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections



16.22 GB, of which 16.11 GB is the frozen weights


- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections




- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay
- The adapter file is 13.6 MB, against a 16.1 GB base model

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections



- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay
- The adapter file is 13.6 MB, against a 16.1 GB base model
- What needed more than one accelerator this morning now needs one



Section 4

What the low rank assumption costs

4

Low rank is a hypothesis, not a theorem



The rank LoRA imposes

- Hu et al. wrote “we hypothesize the updates to the weights also have a low intrinsic rank”

Low rank is a hypothesis, not a theorem



The rank LoRA imposes

- Hu et al. wrote “we hypothesize the updates to the weights also have a low intrinsic rank”
- Aghajanyan measured a random subspace, not a weight update: $d_{90} = 207$ for RoBERTa-Large on MRPC

1

The rank measured in full fine-tuning’s own update



- [Hu et al.](#) wrote “we hypothesize the updates to the weights also have a low intrinsic rank”
- [Aghajanyan](#) measured a random subspace, not a weight update: $d_{90} = 207$ for RoBERTa-Large on MRPC
- [Biderman](#) measured the update itself: 10 to 100 times typical LoRA ranks, and higher in the MLP than in attention

1

The rank measured in full fine-tuning’s own update



- [Hu et al.](#) wrote “we hypothesize the updates to the weights also have a low intrinsic rank”
- [Aghajanyan](#) measured a random subspace, not a weight update: $d_{90} = 207$ for RoBERTa-Large on MRPC
- [Biderman](#) measured the update itself: 10 to 100 times typical LoRA ranks, and higher in the MLP than in attention
- Low rank is a constraint you impose. It is often harmless, and nobody proved it has to hold


Three comparisons with full fine-tuning

Code instruction tuning
HumanEval




Three comparisons with full fine-tuning


Code instruction tuning
HumanEval




Code continued pretraining
HumanEval



Three comparisons with full fine-tuning




Three comparisons with full fine-tuning




- Tulu 2, eight evaluations averaged: 47.7 against 53.0 at 7B, and 71.0 against 73.4 at 70B

Three comparisons with full fine-tuning




- Tulu 2, eight evaluations averaged: 47.7 against 53.0 at 7B, and 71.0 against 73.4 at 70B
- The gap is worst on open ended generation, and it shrinks with model size

Three comparisons with full fine-tuning




- Tulu 2, eight evaluations averaged: 47.7 against 53.0 at 7B, and 71.0 against 73.4 at 70B
- The gap is worst on open ended generation, and it shrinks with model size
- Higher is better in all three panels, so the third is the one LoRA wins

Rank and target modules are real knobs




q_proj
k_proj
v_proj
o_proj
mlp_up
mlp_down


Rank and target modules are real knobs




HumanEval, code instruction tuning




Rank and target modules are real knobs




HumanEval, code instruction tuning




Rank and target modules are real knobs



HumanEval, code instruction tuning




Rank and target modules are real knobs




HumanEval, code instruction tuning

PEFT's default




- `q_proj` and `v_proj` is PEFT's default, faithfully reproducing the 2021 paper

Rank and target modules are real knobs




HumanEval, code instruction tuning




What every study since recommends

- `q_proj` and `v_proj` is PEFT's default, faithfully reproducing the 2021 paper
- QLoRA: adapters on all linear layers are required to match full fine-tuning

Rank and target modules are real knobs



HumanEval, code instruction tuning



What every study since recommends

- `q_proj` and `v_proj` is PEFT's default, faithfully reproducing the 2021 paper
- **QLoRA**: adapters on all linear layers are required to match full fine-tuning
- **Biderman**: most of the gain comes from the MLP, not from attention




Section 5


One base model, many adapters

5


Two thousand adapters on one machine




Two thousand adapters on one machine




Two thousand adapters on one machine



Two thousand adapters on one machine




Two thousand adapters on one machine




- 100 adapted copies of GPT-3: about 354 GB, against about 35 TB

Two thousand adapters on one machine




- 100 adapted copies of GPT-3: about 354 GB, against about 35 TB
- S-LoRA on one A100: 8.05 requests per second at 5 adapters, 7.61 at 2000

Two thousand adapters on one machine



- 100 adapted copies of GPT-3: about 354 GB, against about 35 TB
- S-LoRA on one A100: 8.05 requests per second at 5 adapters, 7.61 at 2000
- The packed baseline manages 2.04 at 5 adapters and runs out of memory at 100

Two thousand adapters on one machine



- 100 adapted copies of GPT-3: about 354 GB, against about 35 TB
- S-LoRA on one A100: 8.05 requests per second at 5 adapters, 7.61 at 2000
- The packed baseline manages 2.04 at 5 adapters and runs out of memory at 100

The catch: merging $W_0 + BA$ forces one copy of the weights per adapter, so you get the merge or the shared batch.

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Prefix and
prompt tuning

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Prefix and
prompt tuning

Preference
tuning and DPO

Preference data is
Weekend 2's material

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Prefix and
prompt tuning

Preference
tuning and DPO

Preference data is
Weekend 2's material

- Thirty minutes buys one mechanism, priced with the same ledger as the rest of the weekend

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Prefix and
prompt tuning

Preference
tuning and DPO

Preference data is
Weekend 2's material

- Thirty minutes buys one mechanism, priced with the same ledger as the rest of the weekend
- LoRA is not a cheaper full fine-tune. It is a different point on a trade off, and often the point you want

Four things this half hour did not cover

Full fine-tuning
mechanics

Priced by this morning's ledger

QLoRA: the
base at 4 bits

Apertus-8B: 16.1 GB
becomes 4.15 GB

Prefix and
prompt tuning

Preference
tuning and DPO

Preference data is
Weekend 2's material

- Thirty minutes buys one mechanism, priced with the same ledger as the rest of the weekend
- LoRA is not a cheaper full fine-tune. It is a different point on a trade off, and often the point you want
- The 11:00 exercise is where you run it