



Fine-tuning methods for LLMs

Adapting a model you did not train

Carlos Cotrini



Section 1

From pretraining to adapting



You will start from a released model

Pretraining Apertus-70B

6.30×10^{24} FLOP

- Two days have priced one pretraining run

You will start from a released model

Pretraining Apertus-70B

6.30×10^{24} FLOP

Over 6 million
GPU hours

- Two days have priced one pretraining run

You will start from a released model

Pretraining Apertus-70B


6.30×10^{24} FLOP

Over 6 million
GPU hours

About 4.95 GWh


- Two days have priced one pretraining run

You will start from a released model




- Two days have priced one pretraining run
- None of it is a decision you will ever make

You will start from a released model




- Two days have priced one pretraining run
- None of it is a decision you will ever make
- You cannot pretrain, so you **adapt**, and the same ledger prices that

Sixteen bytes per parameter




Sixteen bytes per parameter



- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$


Sixteen bytes per parameter



- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$
- Apertus-8B needs 128.8 GB of state, and a GH200 holds 96 GB

Sixteen bytes per parameter

14 of the 16 exist only for the parameters you train



- Mixed precision Adam, from this morning: $2 + 2 + 4 + 4 + 4$
- Apertus-8B needs 128.8 GB of state, and a GH200 holds 96 GB
- Freeze a parameter and 14 of its 16 bytes disappear




Section 2

One matrix, and a low rank update

2

Twenty-four numbers, or ten


W : 6 outputs $\times$ 4 inputs



All 24 are updated

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs



All 24 are updated

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

A : 1×4

a_1	a_2	a_3	a_4
-------	-------	-------	-------

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

- Full fine-tuning updates all 24

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6 $\times$ 1

b_1
b_2
b_3
b_4
b_5
b_6

A : 1 $\times$ 4

a_1	a_2	a_3	a_4
-------	-------	-------	-------

$b_1 a_1$	$b_1 a_2$	$b_1 a_3$	$b_1 a_4$
$b_2 a_1$	$b_2 a_2$	$b_2 a_3$	$b_2 a_4$
$b_3 a_1$	$b_3 a_2$	$b_3 a_3$	$b_3 a_4$
$b_4 a_1$	$b_4 a_2$	$b_4 a_3$	$b_4 a_4$
$b_5 a_1$	$b_5 a_2$	$b_5 a_3$	$b_5 a_4$
$b_6 a_1$	$b_6 a_2$	$b_6 a_3$	$b_6 a_4$

$$\Delta W = BA$$

- Full fine-tuning updates all 24
- LoRA freezes W and learns B and A : $6 + 4 = 10$ numbers, a ratio of 2.4

Twenty-four numbers, or ten

W : 6 outputs $\times$ 4 inputs

Frozen: no gradient, no moments

B : 6×1

b_1
b_2
b_3
b_4
b_5
b_6

A : 1×4


a_1	a_2	a_3	a_4
-------	-------	-------	-------

$b_1 a_1$	$b_1 a_2$	$b_1 a_3$	$b_1 a_4$
$b_2 a_1$	$b_2 a_2$	$b_2 a_3$	$b_2 a_4$
$b_3 a_1$	$b_3 a_2$	$b_3 a_3$	$b_3 a_4$
$b_4 a_1$	$b_4 a_2$	$b_4 a_3$	$b_4 a_4$
$b_5 a_1$	$b_5 a_2$	$b_5 a_3$	$b_5 a_4$
$b_6 a_1$	$b_6 a_2$	$b_6 a_3$	$b_6 a_4$

$$\Delta W = BA$$


- Full fine-tuning updates all 24
- LoRA freezes W and learns B and A : $6 + 4 = 10$ numbers, a ratio of 2.4
- Every row of BA is the row A , rescaled. That is what rank one means

The adapter is added, not substituted




$$h = W_0 x$$

The adapter is added, not substituted



$$h = W_0x + BAx$$


The adapter is added, not substituted



$$h = W_0x + BAx$$

- Both paths read the same x , and their outputs are added


The adapter is added, not substituted



$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added


The adapter is added, not substituted



$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added
- α/r rescales the adapter, and it behaves like a learning rate

The adapter is added, not substituted




$$h = \underbrace{W_0 x}_{\text{frozen}} + \underbrace{\frac{\alpha}{r} B A x}_{\text{trained}}$$

- Both paths read the same x , and their outputs are added
- α/r rescales the adapter, and it behaves like a learning rate
- The paper's own tables use ratios of 1, 2 and 8; the PEFT default is $r = 8, \alpha = 8$

B starts at zero, A starts random

$$B = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

B starts at zero, A starts random



B starts at zero, A starts random

$B = 0$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

- The adapter does nothing at all on step one

B starts at zero, A starts random

$B = 0$

0
0
0
0
0
0

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$BA = 0$

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one

B starts at zero, A starts random

$B = 0$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one
- So training starts where pretraining stopped, not somewhere random

B starts at zero, A starts random

$B = 0$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$BA = 0$

A random

a	a	a	a
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$$W_0 + BA = W_0$$

The first forward pass is exactly the pretrained model

- The adapter does nothing at all on step one
- So training starts where pretraining stopped, not somewhere random
- If both started at zero, both gradients would be zero and nothing would ever move


Rank three costs more than the matrix

$r = 1$  10


The matrix itself: 24

⋮


Rank three costs more than the matrix




Rank three costs more than the matrix



Rank three costs more than the matrix




Rank three costs more than the matrix




- Break even at $r = dk/(d + k) = 2.4$ on this matrix

Rank three costs more than the matrix



- Break even at $r = dk/(d + k) = 2.4$ on this matrix
- On a square $d \times d$ matrix that is $r = d/2$, and the saving is $d/(2r)$

Rank three costs more than the matrix



- Break even at $r = dk/(d + k) = 2.4$ on this matrix
- On a square $d \times d$ matrix that is $r = d/2$, and the saving is $d/(2r)$
- Llama-2-7B's query projection is 4096×4096 : at $r = 8$ that is 65 536 numbers against 16 777 216, a factor of 256




Section 3

The same picture at real scale

3

☰ The same picture at real scale

Rank 16, and 0.08 percent of the model trains



8.05×10^9 parameters

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
<code>q_proj</code> , <code>v_proj</code> (the PEFT default)	16	6 815 744	0.0846

- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

☰ The same picture at real scale

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
<code>q_proj</code> , <code>v_proj</code> (the PEFT default)	16	6 815 744	0.0846
All six linear layers	16	39 845 888	0.4948

- The lower bar is drawn to scale, which is why you cannot see it

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

Rank 16, and 0.08 percent of the model trains



Target modules	Rank	Trainable	Percent
<code>q_proj</code> , <code>v_proj</code> (the PEFT default)	16	6 815 744	0.0846
All six linear layers	16	39 845 888	0.4948
All six linear layers	64	159 383 552	1.9792
All six linear layers	256	637 534 208	7.9167

- The lower bar is drawn to scale, which is why you cannot see it
- Even rank 256 on all six linear layers stays under 8 percent

Apertus-8B's shape is reconstructed from its published totals: 32 layers, $d_{\text{model}} = 4096$, 8 key and value heads.

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state



☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

Weights and gradients	Optimiser state: 96.6 GB
--------------------------	--------------------------


- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB




- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections




- Optimiser state: 96.6 GB becomes 0.082 GB

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections



16.22 GB, of which 16.11 GB is the frozen weights


- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay

☰ The same picture at real scale

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections




- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay
- The adapter file is 13.6 MB, against a 16.1 GB base model

128.8 GB does not fit; 16.22 GB does

Full fine-tuning: 128.8 GB of state

A GH200 holds 96 GB



LoRA, $r = 16$ on the query and value projections



- Optimiser state: 96.6 GB becomes 0.082 GB
- Total state falls only 7.9 times, because the frozen weights stay
- The adapter file is 13.6 MB, against a 16.1 GB base model
- What needed more than one accelerator this morning now needs one